

Assignment 10

Hendrik Werner s4549775

June 10, 2017

1 Abstract

In this paper I present an easy to understand and follow approach to detect an object's position, and the direction it is facing, from visual input using OpenCV. This is done as a case study of an autonomous football playing robot.

2 Introduction

For a university project Timo Schrijvers and myself wanted to detect the position of a robot and which way it was facing, as well as the position of a ball, using visual input. After looking into different computer vision frameworks, we eventually chose OpenCV. We expected there to be some prebuilt functionality for something as mundane as detecting an arrow but that was not the case.

In this article I present an approach to detect the features we needed; and how this was applied in a robotics project. This article is positioned as an introduction into image detection and does not assume familiarity with the subject. Concepts and terminology are introduced as needed; it serves as a quick overview of image detection and OpenCV in particular.

3 Problem

We built a little robot out of Lego Mindstorm EV3. The robot is really basic and can just drive around, steer, and pick up a ball.

As input for our robot's decision making algorithm we need the position of the ball, as well as the position of the robot, and the direction it faces. The ball is round, so direction does not really apply to it.

For this purpose we mounted a camera above the playing field facing straight down. The ball is a bright yellow, and on the robot we mounted a bright red arrow, pointing forward.

In summary, the problem we need to solve is the following: Given visual input with unique, and brightly colored objects, how can the positions and directions of the objects be extracted?

4 Overview of Computer Vision

Computer Vision is the field of study that deals with the extraction of high level features from digital visual data. In our case this data is the video stream from the camera, and the features are the positions and directions of the ball and car.

There are many applications for Computer Vision, most of which are concerned with automation of tasks previously done by humans. In our case playing football, though there are of course more practical applications as well.

Computer Vision is used by huge, international companies to generate billions in revenue [7]. Tesla, for example, heavily invests in Computer Vision for its application in self driving vehicle technology [14] [15].

5 Overview of OpenCV

OpenCV (Open Computer Vision) is a cross-platform, open source computer vision library. Development was started by Intel in 1999 as a research project.

It is released under the BSD license which allows for commercial and non-commercial free use [5].

Today it is one of the industry standards for Computer Vision, used and sponsored by companies such as Google, Microsoft, and Intel, among many others. It has interfaces for C, C++, Python, Java, and MATLAB [2].

OpenCV is highly optimized and geared towards real time systems [2]. This is important for many applications in automation. For example, your car must be able to process visual information in real time to avoid collisions. For our purposes this is not of great importance but it is nice to be able to visualize the algorithms working. Quicker response times are also generally nice.

The OpenCV API provides a broad range of features, a list of which can be found on their website [11]. This includes some non-free components [12], as well as deprecated modules and experimental / contributed functionality.

6 Solution

6.1 Preprocessing

OpenCV reads the image in as BGR (blue, green, red), where a color's hue is split across 3 different channels. We want to filter on the hue, so it is the easiest solution to convert the color to HSV (hue, saturation, value), which packs hue information into a single channel.

```
image_hsv = cv2.cvtColor(image, cv2.COLOR_BGR2HSV)
```

The images we capture from the camera are pretty noisy, so we apply a median blur. A bilateral filter is even better because it reduces noise while retaining edges. For our purposes a median blur is sufficient, so we use it because its performance is better. [13]

```
image_hsv = cv2.medianBlur(image_hsv, 3)
```

We made a generic object finder class, that does the following things:

1. Make a mask. This is a black and white image of the same dimensions as the input image, in which all pixels falling within a certain color

range are white, and the rest is black. It basically maps all pixels to whether or not they are in the color range.

We use OpenCV's *inRange* function for this:

```
mask = cv2.inRange(  
    image_hsv ,  
    threshold [ "lower" ] ,  
    threshold [ "upper" ] ,  
)
```

2. Find the contours. The contours are the outlines of all the shapes in our mask.

We use *findContours*:

```
contours = cv2.findContours(  
    mask ,  
    cv2.RETR_LIST ,  
    cv2.CHAIN_APPROX_SIMPLE ,  
)
```

3. Filter out all contours with an area below some threshold:

```
contours = [  
    c for c in contours  
    if cv2.contourArea(c) >= minimum_area  
]
```

6.2 Finding the ball's position

This is the easy part of our project. To find the ball we just need to find the center of the minimum enclosing circle of the biggest blob of pixels that have the right color. We assume that the general steps described above have been done already.

contourArea and *minEnclosingCircle* can be used for this:

```
contour = max(self.contours , key=cv2.contourArea)  
pos , radius = cv2.minEnclosingCircle(contour)
```

6.3 Finding the robot's position and direction

One could use the ball finding technique for the robot too, but that would only tell one it's position. Since we also need the direction it is facing we strapped an arrow to the robot. We look for the center of this arrow, and for the pointy end. With these two points we have our robot vector.

After applying the general steps described above, we take all contours and approximate them using *approxPolyDP*:

```
contour = cv2.approxPolyDP(contour, 10, closed=True).  
    squeeze()
```

Next we extract all the lines and turn them into vector combinations:

```
def angle_heuristic(  
    desired_angle: float,  
    v1: np.ndarray,  
    v2: np.ndarray,  
    min_angle: int = 10,  
) -> float:  
    min_angle = as_rad(min_angle)  
    inner_angle = angle(v1, v2, True)  
    if inner_angle < min_angle:  
        return np.pi * 2  
    return min(  
        abs(desired_angle - inner_angle),  
        abs(desired_angle - np.pi + inner_angle),  
    )  
  
vectors = [  
    (i, [x2 - x1, y2 - y1])  
    for i, (x1, y1, x2, y2) in enumerate(lines)  
]  
combinations = [  
    (angle_heuristic(arrow_angle, c[0][1], c[1][1]), c)  
    for c in itertools.combinations(vectors, 2)  
]
```

Then we add the best combination we found to our candidates list. After we have done this for all contours, we select the best of all of our candidates,

which is a pair of lines representing the best arrow we could find.

The intersection of these two lines is the tip of the arrow, while the end of the shaft is the center of all the points that make up the two lines:

```
arr_pos = cv2.minEnclosingCircle(  
    np.append(arr_lines[0], arr_lines[1]).reshape((4,  
        2))  
)[0]
```

7 Application

We use the information we extracted as input for an algorithm that decides on a plan of action, that describes what our robot is supposed to do. This information is sent to an executor, which controls the robot via Bluetooth.

Even though the techniques were basic, and we extracted very limited information from the camera input, this was enough to make an arbitrarily placed robot find an arbitrarily placed ball, as long as they are visible to the camera.

This demonstrates the ease with which Computer Vision can be deployed even by amateurs, with the help of frameworks such as OpenCV.

8 Conclusion

OpenCV is a great tool for Computer Vision, which provides a basis to build sophisticated feature extraction applications using visual input. These features can in turn be used for automation in many different areas.

The documentation is not very comprehensive, especially for non-native interfaces of OpenCV. This can make it unnecessarily hard for beginners to get into the matter. With the right reading material however, it can be easy to quickly achieve great results, even for beginners; by building on the solid foundation of the OpenCV API.

The application for Computer Vision are diverse and it can be used from

little hobby or school projects, to building the next generation of self driving vehicles.

9 Further Reading

- The full source code, the planner, the executor, the specifications for building the robot, and additional documentation can be found on the project's GitHub page[8].
- The OpenCV 3 Cookbook [10] is a complete introduction into the third version of the OpenCV library. "You will be presented with a variety of computer vision algorithms and exposed to important concepts in image and video analysis that will enable you to build your own computer vision applications." [4]
- If you have Python experience you might want to look at OpenCV Computer Vision with Python [9], which focuses on Python.

"This book has practical, project-based tutorials for Python developers and hobbyists who want to get started with computer vision with OpenCV and Python. It is a hands-on guide that covers the fundamental tasks of computer vision, capturing, filtering, and analyzing images, with step-by-step instructions for writing both an application and reusable library classes." [1]

- If you are more of a C++ fan, or you want to use OpenCV through its native interface, maybe you are interested in Learning OpenCV 3 Application Development [6].

"This book provides the steps to build and deploy an end-to-end application in the domain of computer vision using OpenCV/C++.". Topics covered are: how images are stored and processed by OpenCV, OpenCV-specific jargon, image traversal and pixel-wise operations, filtering, thresholding, edge detection, face detection, and much more. [3]

Keep in mind that the C++ interface has much better documentation because of being the native language of OpenCV. However, generally speaking, Python is used for a lot of automation and Machine Learning, and provides a wealth of other supporting libraries for a wide variety of tasks. Python is

also considered to be more concise than C++, which is paid for in reduced performance.

Do not focus too much on the choice of languages though, not even the choice of frameworks. The general concepts, algorithms, and approaches are much more generic. They will serve you well no matter the language or framework you happen to be working in.

References

- [1] *About Computer Vision with Python*. URL: <https://www.packtpub.com/application-development/opencv-computer-vision-python>.
- [2] *About OpenCV*. URL: <http://opencv.org/about.html>.
- [3] *About OpenCV 3 Application Development*. URL: <https://www.packtpub.com/application-development/learning-opencv-3-application-development>.
- [4] *About OpenCV 3 Cookbook - Third Edition*. URL: <https://www.packtpub.com/application-development/opencv-3-computer-vision-application-programming-cookbook-third-edition>.
- [5] Gary Bradski and Adrian Keahler. *Learning OpenCV*. O'Reilly, 2008.
- [6] Samyak Datta. *Learning OpenCV 3 Application Development*. Packt, 2016.
- [7] *Elon Musk's Net Worth Rises To All-Time High*. URL: <https://www.forbes.com/sites/noahkirsch/2017/04/08/as-tesla-tops-ford-elon-musks-net-worth-rises-1-8-billion-to-all-time-high>.
- [8] Timo Schrijvers Hendrik Werner. *Football Project - SoccerBot*. URL: https://github.com/NewDevices/football_project/tree/subprocess.
- [9] Joseph Howse. *OpenCV Computer Vision with Python*. Packt, 2013.
- [10] Robert Laganriere. *OpenCV 3 Computer Vision Application Programming Cookbook - Third Edition*. Packt, 2017.
- [11] *OpenCV API Reference*. URL: <http://docs.opencv.org/3.0-beta/modules/refman.html>.
- [12] *OpenCV nonfree - Feature Detection and Description*. URL: http://docs.opencv.org/2.4/modules/nonfree/doc/feature_detection.html.

- [13] *Smoothing Images*. URL: http://docs.opencv.org/3.0-beta/doc/tutorials/imgproc/gaussian_median_blur_bilateral_filter/gaussian_median_blur_bilateral_filter.html.
- [14] *Tesla Autopilot*. URL: <https://www.tesla.com/autopilot>.
- [15] *Tesla is about to increase its lead in semi-autonomous driving w/ 'Tesla Vision': computer vision based on NVIDIA's parallel computing*. URL: <https://electrek.co/2016/10/10/tesla-vision-autopilot-autonomous-driving-nvidia-cuda/>.

10 Feedback taken into account

- Some minor grammatical errors.
- Use "I" and "we" in the correct places.
"I" wrote the article but "we" did the research as a team.
- Too many verbs.

before

In this article I want to present the approach I took to detect the features I needed.

after

In this article I present an approach to detect the features we needed.

- Sentence is too vague. What does "disappointed" actually mean?

before

I expected there to be some prebuilt functionality for something as mundane as detecting an arrow but was disappointed.

after

We expected there to be some prebuilt functionality for something as mundane as detecting an arrow but that was not the case.

- Do not use informal forms like "isn't".
- Reference a list of features instead of enumerating them.

before

The OpenCV API provides facilities for image (pre)processing, persistence, clustering, high-level GUI, video analysis, camera calibration, 2d-feature extraction, object detection, Machine Learning, GPU acceleration, computational photography, image stitching, OpenCL acceleration, super resolution, 3d-visualization and even some non-free components [12].

In addition to that there are some deprecated modules and experimental / contributed functionality. [11]

after

The OpenCV API provides a broad range of features, a list of which can be found on their website [11]. This includes some non-free components [12], as well as deprecated modules and experimental / contributed functionality.

- Finish half sentence.

before

OpenCV is a great tool for Computer Vision, which provides a basis to build sophisticated

after

OpenCV is a great tool for Computer Vision, which provides a basis to build sophisticated feature extraction applications using visual input. These features can in turn be used for automation in many different areas.

- Clarify why Tesla is brought up.

before

Tesla heavily invests in Computer Vision for its application in self driving vehicle technology [14] [15].

after

Computer Vision is used by huge, international companies to generate billions in revenue [7]. Tesla, for example, heavily invests in Computer Vision for its application in self driving vehicle technology [14] [15].

- Clearly identify the sentence as an example in the beginning.

before

Your car must be able to process visual information in real time to avoid collisions, for example.

after

For example, your car must be able to process visual information in real time to avoid collisions.

- Explain what "the robot" is.

before

after

We built a little robot out of Lego Mindstorm EV3. The robot is really basic and can just drive around, steer, and pick up a ball.